

Classify different data sets

Basic includes

In [2]:

```
1  ▾ # Using pandas to load the csv file
2  import pandas as pd
3
4  import numpy as np
5  import matplotlib.pyplot as plt
6
7  from keras import models
8  from keras import layers
9  from keras import callbacks
10 from keras.utils import to_categorical
11
12 # reuters and fashin mnist data set from keras
13 from keras.datasets import reuters
14 from keras.datasets import fashion_mnist
15
16 # needed to preprocess text
17 from keras.preprocessing.text import Tokenizer
```

executed in 3.88s, finished 13:44:52 2019-02-19

Using TensorFlow backend.

Classify the Fashion Mnist


In [86]:

```
1 (fashion_train_data, fashion_train_labels), (fashion_test_data, fashion_test_l
2 ▼ fashion_class_labels = [
3     'T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Shirt',
4 ]
5 print(fashion_train_data.shape)
6
7 test_index = 10
8
9 plt.title("Label: " + fashion_class_labels[fashion_train_labels[test_index]])
10 plt.imshow(fashion_train_data[test_index], cmap="gray")
```

(60000, 28, 28)

Out[86]:

<matplotlib.image.AxesImage at 0x7facef2f0940>



TO DO: Preprocess the data

1. Normalize the input data set
2. Perform one hot encoding
3. Create a train, test, and validation set

In [0]:

```
1  ▼  # Normalize the input data set
2      # flatten images
3      fashion_train_data = fashion_train_data.reshape((60000, 784))
4      fashion_train_data = fashion_train_data.astype('float32') / 255
5
6      fashion_test_data = fashion_test_data.reshape((10000, 784))
7      fashion_test_data = fashion_test_data.astype('float32') / 255
8
9      # one hot encoding
10     fashion_train_labels = to_categorical(fashion_train_labels)
11     fashion_test_labels = to_categorical(fashion_test_labels)
12
13     validation_set_labels = fashion_train_labels[50000:]
14     validation_set = fashion_train_data[50000:]
15
16     training_set_labels = fashion_train_labels[:50000]
17     training_set = fashion_train_data[:50000]
```

TO DO: Define and train a network, then plot the accuracy of the training, validation, and testing

1. Use a validation set
2. Propose and train a network
3. Print the history of the training
4. Evaluate with a test set

In [88]:

```
1  # Crate NN
2
3  nn_fashion_model = models.Sequential()
4  fashion_dropout = 0.3
5
6  nn_fashion_model.add(layers.Dense(1024, activation= "relu", input_shape= (784,
7
8  nn_fashion_model.add(layers.Dropout(fashion_dropout))
9
10 nn_fashion_model.add(layers.Dense(256, activation="relu"))
11 nn_fashion_model.add(layers.Dense(128, activation="relu"))
12
13 nn_fashion_model.add(layers.Dropout(fashion_dropout))
14
15 # Last layer, same size has the number of categories
16 nn_fashion_model.add(layers.Dense(10, activation="softmax"))
17
18
19
20 nn_fashion_early_stops = [
21     callbacks.EarlyStopping(monitor= 'val_loss', patience= 4)
22 ]
23
24 nn_fashion_model.compile(
25     loss= "categorical_crossentropy", optimizer= "adam", metrics= ["accuracy"]
26 )
27
28 nn_fashion_model.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_47 (Dense)	(None, 1024)	803840
<hr/>		
dropout_17 (Dropout)	(None, 1024)	0
<hr/>		
dense_48 (Dense)	(None, 256)	262400
<hr/>		
dense_49 (Dense)	(None, 128)	32896
<hr/>		
dropout_18 (Dropout)	(None, 128)	0
<hr/>		
dense_50 (Dense)	(None, 10)	1290
=====		
Total params: 1,100,426		
Trainable params: 1,100,426		
Non-trainable params: 0		
<hr/>		

In [89]:

```
1  ▼ # Train the NN model
2    fashion_epochs = 16
3  ▼ nn_fashion_history = nn_fashion_model.fit(
4      fashion_train_data,
5      fashion_train_labels,
6      batch_size= 1024,
7      epochs= fashion_epochs,
8      verbose= 2,
9      callbacks= nn_fashion_early_stops,
10     validation_data= (validation_set, validation_set_labels)
11 )
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/16

- 8s - loss: 0.7594 - acc: 0.7338 - val_loss: 0.4379 - val_acc: 0.843

4

Epoch 2/16

- 7s - loss: 0.4523 - acc: 0.8404 - val_loss: 0.3780 - val_acc: 0.862

7

Epoch 3/16

- 7s - loss: 0.3956 - acc: 0.8588 - val_loss: 0.3368 - val_acc: 0.877

7

Epoch 4/16

- 7s - loss: 0.3634 - acc: 0.8692 - val_loss: 0.3245 - val_acc: 0.881

2

Epoch 5/16

- 7s - loss: 0.3418 - acc: 0.8758 - val_loss: 0.3016 - val_acc: 0.889

5

Epoch 6/16

- 8s - loss: 0.3206 - acc: 0.8832 - val_loss: 0.2827 - val_acc: 0.895

7

Epoch 7/16

- 8s - loss: 0.3106 - acc: 0.8865 - val_loss: 0.2850 - val_acc: 0.894

9

Epoch 8/16

- 7s - loss: 0.3032 - acc: 0.8897 - val_loss: 0.2718 - val_acc: 0.894

9

Epoch 9/16

- 7s - loss: 0.2951 - acc: 0.8915 - val_loss: 0.2719 - val_acc: 0.897

2

Epoch 10/16

- 8s - loss: 0.2813 - acc: 0.8959 - val_loss: 0.2448 - val_acc: 0.910

5

Epoch 11/16

- 7s - loss: 0.2728 - acc: 0.8990 - val_loss: 0.2459 - val_acc: 0.905

0

Epoch 12/16

- 7s - loss: 0.2696 - acc: 0.8997 - val_loss: 0.2335 - val_acc: 0.910

1

Epoch 13/16
- 7s - loss: 0.2591 - acc: 0.9043 - val_loss: 0.2241 - val_acc: 0.9148
Epoch 14/16
- 7s - loss: 0.2564 - acc: 0.9056 - val_loss: 0.2170 - val_acc: 0.9170
Epoch 15/16
- 7s - loss: 0.2504 - acc: 0.9066 - val_loss: 0.2166 - val_acc: 0.9173
Epoch 16/16
- 7s - loss: 0.2433 - acc: 0.9080 - val_loss: 0.2176 - val_acc: 0.9173

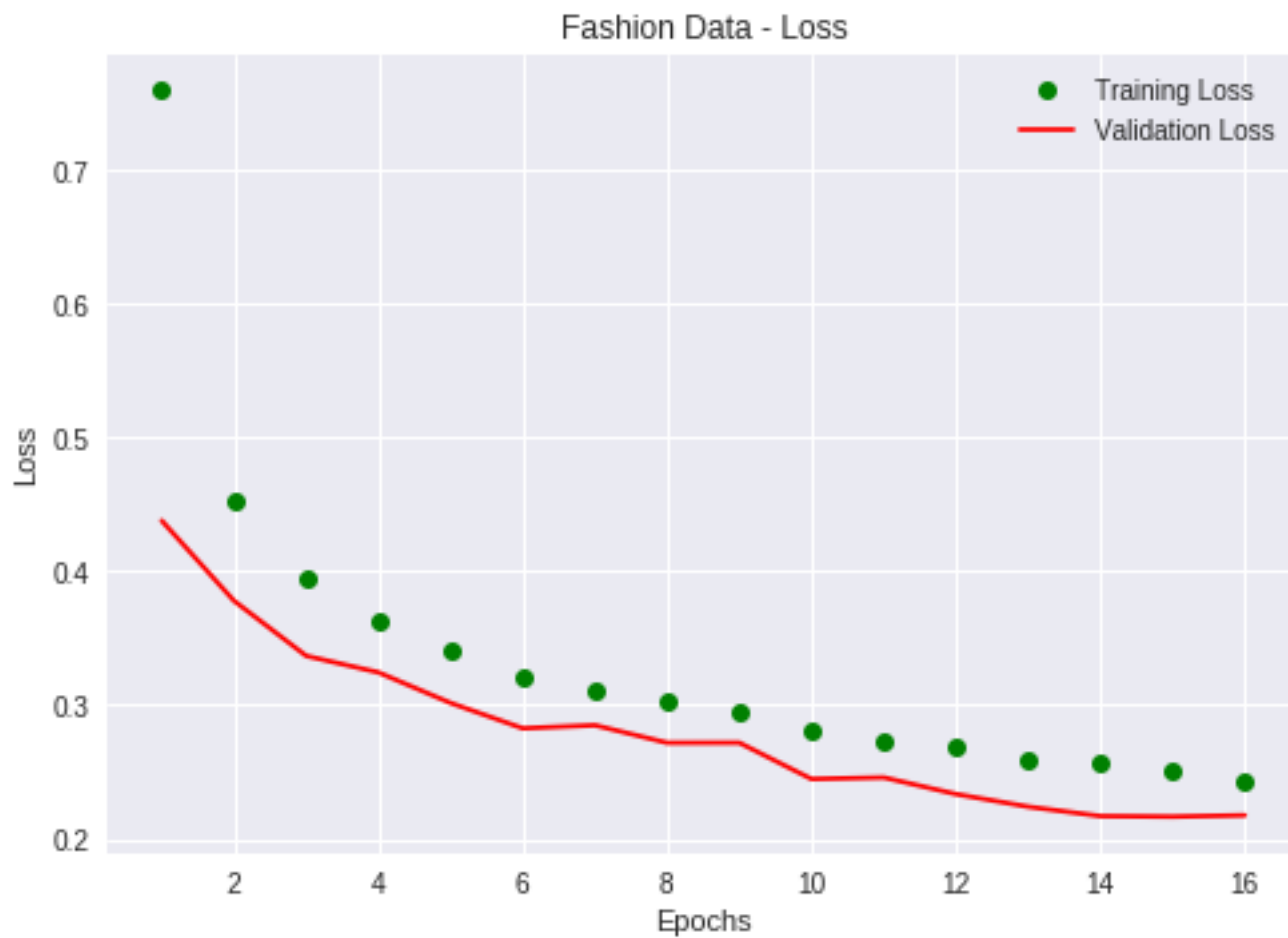
In [90]:

1	fashion_result = nn_fashion_model.evaluate(fashion_test_data, fashion_test_labels)
2	print('Fashion score: {}'.format(fashion_result[1]*100))

10000/10000 [=====] - 1s 105us/step
Fashion score: 88.74%

In [91]:

```
1 fashion_history = nn_fashion_history.history
2 fashion_loss = fashion_history['loss']
3 fashion_val_loss = fashion_history['val_loss']
4 fashion_epochs = range(1, len(fashion_loss) + 1)
5
6 plt.plot(fashion_epochs, fashion_loss, 'go', label='Training Loss')
7 plt.plot(fashion_epochs, fashion_val_loss, 'r', label='Validation Loss')
8
9 plt.title('Fashion Data - Loss')
10 plt.xlabel('Epochs')
11 plt.ylabel('Loss')
12 plt.legend()
13
14 plt.show()
```



In [92]:

```
1 fashion_acc      = fashion_history['acc']
2 fashion_val_acc  = fashion_history['val_acc']
3
4 plt.plot(fashion_epochs, fashion_acc, 'go', label='Training Acc')
5 plt.plot(fashion_epochs, fashion_val_acc, 'r', label='Validation Acc')
6
7 plt.title('Fashion Data - Accuracy')
8 plt.xlabel('Epochs')
9 plt.ylabel('Loss')
10 plt.legend()
11
12 plt.show()
```



Fashion conclusion

In this model I first did over fitting so making the model network complex, from that point I moved some hyper parameters and successfully achieved a score above 85%.

Classifying newswires

Build a network to classify Reuters newswires into 46 different mutually-exclusive topics.

Load and review the data

In [93]:

```
1  reuters_max_words = 10000
2  (reuters_train_data, reuters_train_labels), (reuters_test_data, reuters_test_l
3
4  print(reuters_train_data.shape)
5  print(reuters_train_labels.shape)
6  print(reuters_train_data[0])
7  print(reuters_train_labels[0])
8
9  print(set(reuters_train_labels))
```

(8982,)

(8982,)

```
[1, 2, 2, 8, 43, 10, 447, 5, 25, 207, 270, 5, 3095, 111, 16, 369, 186,
90, 67, 7, 89, 5, 19, 102, 6, 19, 124, 15, 90, 67, 84, 22, 482, 26, 7,
48, 4, 49, 8, 864, 39, 209, 154, 6, 151, 6, 83, 11, 15, 22, 155, 11, 1
5, 7, 48, 9, 4579, 1005, 504, 6, 258, 6, 272, 11, 15, 22, 134, 44, 11,
15, 16, 8, 197, 1245, 90, 67, 52, 29, 209, 30, 32, 132, 6, 109, 15, 17
, 12]
```

3

```
{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19,
20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37
, 38, 39, 40, 41, 42, 43, 44, 45}
```

Load the word index to decode the train data.

In [94]:

```
1 word_index = Reuters.get_word_index()
2
3 reverse_index = dict([(value+3, key) for (key, value) in word_index.items()])
4
5 reverse_index[0] = "<PAD>"
6 reverse_index[1] = "<START>"
7 reverse_index[2] = "<UNKNOWN>" # unknown
8 reverse_index[3] = "<UNUSED>"
9
10 decoded_review = ' '.join([reverse_index.get(i, '?') for i in Reuters_train_data])
11
12 print(decoded_review)
```

<START> <UNKNOWN> <UNKNOWN> said as a result of its december acquisiti
on of space co it expects earnings per share in 1987 of 1 15 to 1 30 d
lrs per share up from 70 cts in 1986 the company said pretax net shoul
d rise to nine to 10 mln dlrs from six mln dlrs in 1986 and rental ope
ration revenues to 19 to 22 mln dlrs from 12 5 mln dlrs it said cash f
low per share this year should be 2 50 to three dlrs reuter 3

TO DO: Preprocess the data

1. Normalize the input data set
2. Perform one hot encoding
3. Create a train, test, and validation set

In [95]:

```
1 tokenizer = Tokenizer(num_words= reuters_max_words)
2
3 ▼ reuters_train_data_token = tokenizer.sequences_to_matrix(
4     reuters_train_data, mode="binary"
5 )
6 ▼ reuters_test_data_token = tokenizer.sequences_to_matrix(
7     reuters_test_data, mode="binary"
8 )
9
10 reuters_one_hot_train_labels = to_categorical(reuters_train_labels)
11 reuters_one_hot_test_labels = to_categorical(reuters_test_labels)
12
13 reuters_val_data = reuters_train_data_token[:1000]
14 reuters_val_labels = reuters_one_hot_train_labels[:1000]
15
16 reuters_train_data = reuters_train_data_token[1000:]
17 reuters_train_labels = reuters_one_hot_train_labels[1000:]
18
19 print('train:')
20 print(reuters_train_data.shape)
21 print(reuters_train_labels.shape)
22 print('val:')
23 print(reuters_val_data.shape)
24 print(reuters_val_labels.shape)
25 print('test:')
26 print(reuters_test_data_token.shape)
27 print(reuters_one_hot_test_labels.shape)
```

```
train:
(7982, 10000)
(7982, 46)
val:
(1000, 10000)
(1000, 46)
test:
(2246, 10000)
(2246, 46)
```

TO DO: Define and train a network, then plot the accuracy of the training, validation, and testing

1. Use a validation set
2. Propose and train a network
3. Print the history of the training
4. Evaluate with a test set

In [96]:

```
1  reuters_model = models.Sequential()
2  reuters_dropout = 0.2
3
4  reuters_model.add(layers.Dense(1024, activation="tanh", input_dim=10000))
5
6  reuters_model.add(layers.Dropout(reuters_dropout))
7
8  reuters_model.add(layers.Dense(256, activation="relu"))
9  reuters_model.add(layers.Dense(128, activation="relu"))
10
11 reuters_model.add(layers.Dropout(reuters_dropout))
12
13 reuters_model.add(layers.Dense(46, activation="softmax"))
14
15 ▼ reuters_model.compile(
16     loss= "categorical_crossentropy",
17     optimizer= "adamax",
18     metrics= ["accuracy"]
19 )
20
21 reuters_model.summary()
22
23
24 ▼ reuters_early_stops = [
25     callbacks.EarlyStopping(monitor= 'val_loss', patience= 4),
26     callbacks.EarlyStopping(monitor= 'val_acc',  patience= 5)
27 ]
```

Layer (type)	Output Shape	Param #
=====		
dense_51 (Dense)	(None, 1024)	10241024
<hr/>		
dropout_19 (Dropout)	(None, 1024)	0
<hr/>		
dense_52 (Dense)	(None, 256)	262400
<hr/>		
dense_53 (Dense)	(None, 128)	32896
<hr/>		
dropout_20 (Dropout)	(None, 128)	0
<hr/>		
dense_54 (Dense)	(None, 46)	5934
=====		
Total params: 10,542,254		
Trainable params: 10,542,254		
Non-trainable params: 0		

In [97]:

```
1  ▼ reuters_model_history = reuters_model.fit(  
2      reuters_train_data,  
3      reuters_train_labels,  
4      batch_size= 1024,  
5      epochs= 16,  
6      verbose= 2,  
7      callbacks= reuters_early_stops,  
8      validation_data= (reuters_val_data, reuters_val_labels)  
9  )
```

Train on 7982 samples, validate on 1000 samples

Epoch 1/16

- 8s - loss: 2.5642 - acc: 0.4818 - val_loss: 1.5615 - val_acc: 0.6560

Epoch 2/16

- 7s - loss: 1.3029 - acc: 0.7129 - val_loss: 1.1722 - val_acc: 0.7310

Epoch 3/16

- 7s - loss: 0.9083 - acc: 0.7968 - val_loss: 0.9980 - val_acc: 0.7890

Epoch 4/16

- 7s - loss: 0.6427 - acc: 0.8513 - val_loss: 0.9031 - val_acc: 0.8120

Epoch 5/16

- 7s - loss: 0.4676 - acc: 0.8968 - val_loss: 0.8801 - val_acc: 0.8220

Epoch 6/16

- 7s - loss: 0.3395 - acc: 0.9266 - val_loss: 0.9141 - val_acc: 0.8120

Epoch 7/16

- 7s - loss: 0.2594 - acc: 0.9381 - val_loss: 0.9113 - val_acc: 0.8210

Epoch 8/16

- 7s - loss: 0.2066 - acc: 0.9461 - val_loss: 0.9337 - val_acc: 0.8160

Epoch 9/16

- 7s - loss: 0.1731 - acc: 0.9544 - val_loss: 0.9800 - val_acc: 0.8180

In [98]:

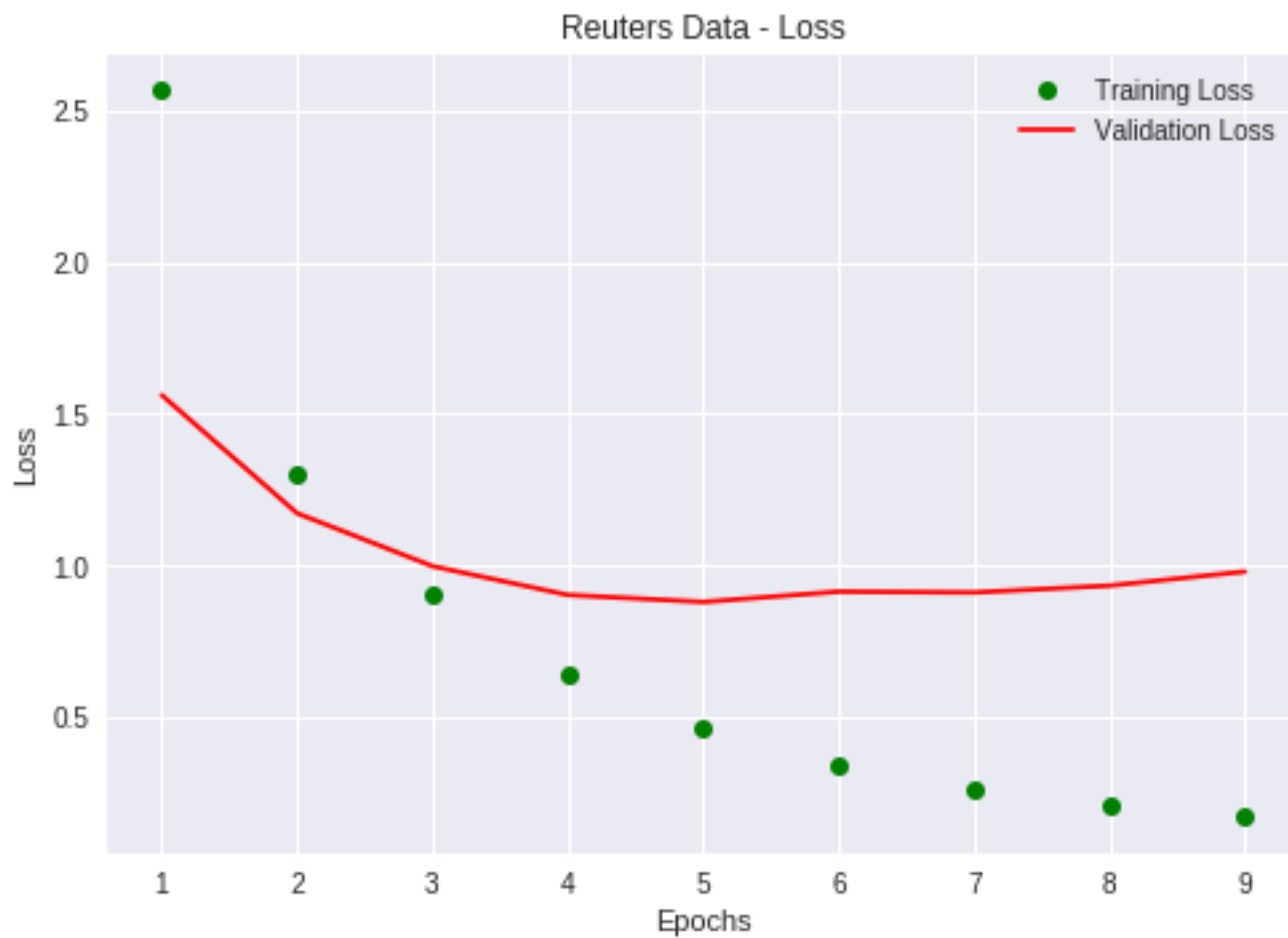
```
1  ▼ reuters_result = reuters_model.evaluate(  
2      reuters_test_data_token,  
3      reuters_one_hot_test_labels  
4  )  
5  
6  print('Fashion score: {}'.format(reuters_result[1]*100))
```

2246/2246 [=====] - 2s 721us/step

Fashion score: 80.00890471950134%

In [100]:

```
1  reuters_history = reuters_model_history.history
2  reuters_loss = reuters_history['loss']
3  reuters_val_loss = reuters_history['val_loss']
4  reuters_epochs = range(1, len(reuters_loss) + 1)
5
6  plt.plot(reuters_epochs, reuters_loss, 'go', label='Training Loss')
7  plt.plot(reuters_epochs, reuters_val_loss, 'r', label='Validation Loss')
8
9  plt.title('Reuters Data - Loss')
10 plt.xlabel('Epochs')
11 plt.ylabel('Loss')
12 plt.legend()
13
14 plt.show()
```



In [101]:

```
1  reuters_acc      = reuters_history['acc']
2  reuters_val_acc  = reuters_history['val_acc']
3
4  plt.plot(reuters_epochs, reuters_acc, 'go', label='Training Acc')
5  plt.plot(reuters_epochs, reuters_val_acc, 'r', label='Validation Acc')
6
7  plt.title('Reuters Data - Accuracy')
8
9  plt.xlabel('Epochs')
10 plt.ylabel('Accuracy')
11 plt.legend()
12 plt.show()
```



Reuters conclusion

1. This model was harder to achieve the accuracy goal. Like the past model, first, I over fitted the model but the accuracy in that point was at it's lowest in all the exercise.
2. Moving hyper parameters wasn't getting me closer to the goal. Neither changing my preprocessing work flow.
3. My final conclusion is that I couldn't get an final accuracy higher than 78% 'cause of the dataset values. My hypothesis is the data is not evenly given to me so I have more cases of 1 category than another.

Predicting Student Admissions

Predict student admissions based on three pieces of data:

- GRE Scores
- GPA Scores
- Class rank

Load and visualize the data


In [3]:

1	<code>student_data = pd.read_csv("student_data.csv")</code>
executed in 25ms, finished 13:44:56 2019-02-19	

Plot of the GRE and the GPA from the data.

In [103]:

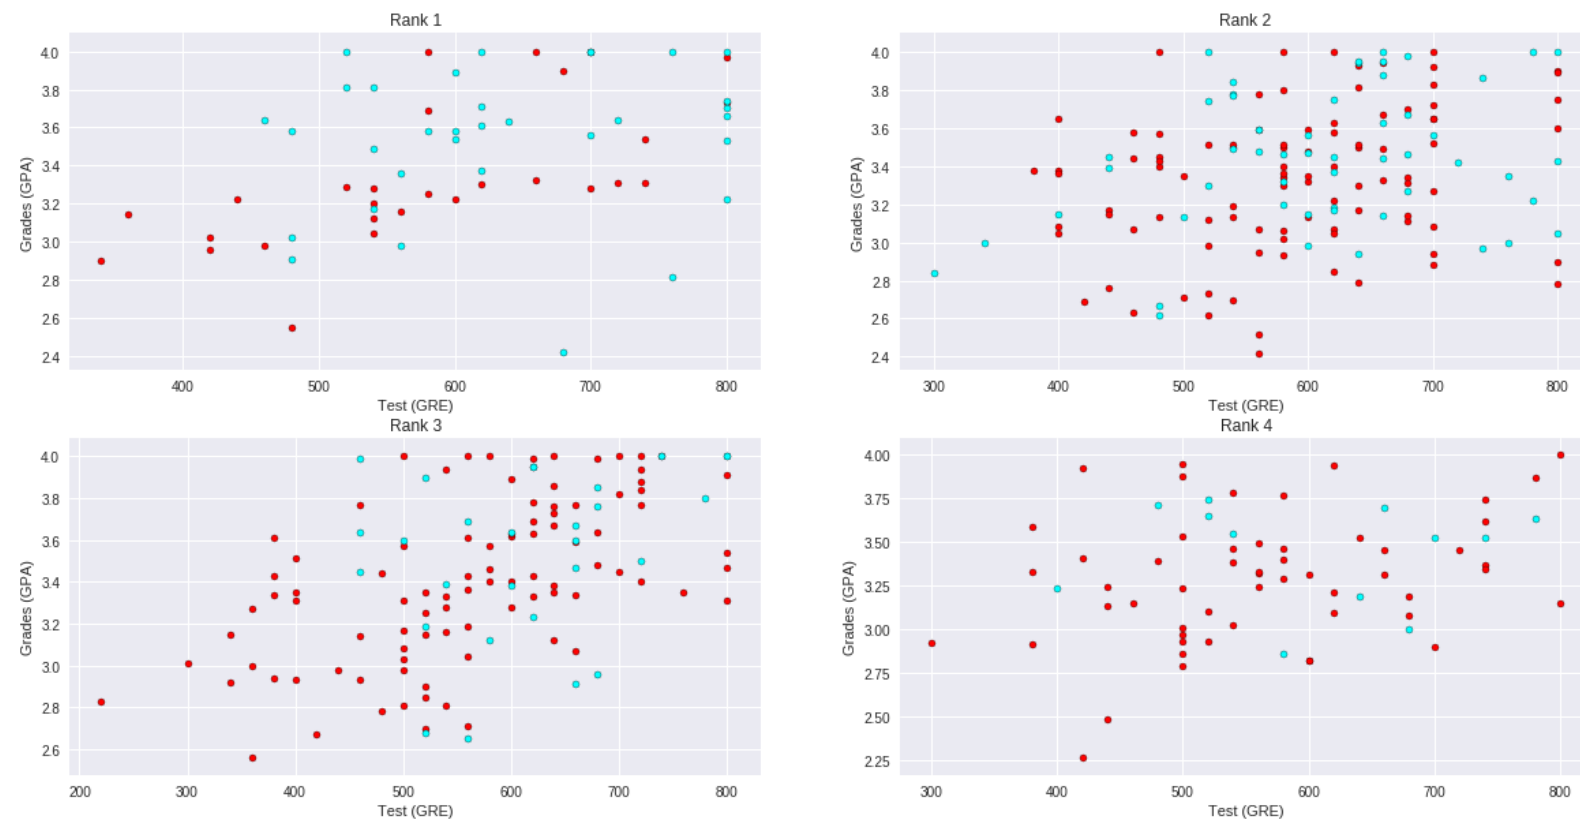
```
1 X = np.array(student_data[["gre", "gpa"]])
2 y = np.array(student_data["admit"])
3 admitted = X[np.argwhere(y==1)]
4 rejected = X[np.argwhere(y==0)]
5 plt.scatter([s[0][0] for s in rejected], [s[0][1] for s in rejected], s = 25,
6 plt.scatter([s[0][0] for s in admitted], [s[0][1] for s in admitted], s = 25,
7 plt.xlabel('Test (GRE)')
8 plt.ylabel('Grades (GPA)')
9
10 plt.show()
```



Plot of the data by class rank.

In [104]:

```
1 f, plots = plt.subplots(2, 2, figsize=(20,10))
2 plots = [plot for sublist in plots for plot in sublist]
3
4 ▼ for idx, plot in enumerate(plots):
5     data_rank = student_data[student_data["rank"]==idx+1]
6     plot.set_title("Rank " + str(idx+1))
7     X = np.array(data_rank[["gre", "gpa"]])
8     y = np.array(data_rank["admit"])
9     admitted = X[np.argwhere(y==1)]
10    rejected = X[np.argwhere(y==0)]
11    plot.scatter([s[0][0] for s in rejected], [s[0][1] for s in rejected], s =
12    plot.scatter([s[0][0] for s in admitted], [s[0][1] for s in admitted], s =
13    plot.set_xlabel('Test (GRE)')
14    plot.set_ylabel('Grades (GPA)')
15
```



TO DO: Preprocess the data

1. Normalize the input data set
2. Perform one hot encoding
3. Create a train, test, and validation set

In [6]:

```
1  # Replace nan with 0
2  student_data.fillna(value= 0, inplace= True)
3  # Shuffle the dataframe with pandas
4  student_data = student_data.sample(frac= 1).reset_index(drop= True)
5
6  # x
7  gre, gpa = np.array(student_data['gre']), np.array(student_data['gpa'])
8  # y
9  admit, rank = np.array(student_data['admit']), np.array(student_data['rank'])
10
11 # Make everything in range from 0 - 1 (Normal distribution)
12 gpa = (gpa - gpa.mean(axis= 0)) / gpa.std(axis= 0)
13 gre = (gre - gre.mean(axis= 0)) / gre.std(axis= 0)
14
15 normalized_student_data = np.zeros((len(gpa), 2))
16 normalized_student_data[:,0], normalized_student_data[:,1] = gpa, gre
17
18 print(normalized_student_data.shape)
19
20 # one hot encoding
21 rank_one_hot = to_categorical(rank)
22
23 # train: 0 300, test: 300 350, val: 350 4000
24 student_train_data = normalized_student_data[:300]
25 student_train_labels = rank_one_hot[:300]
26
27 student_test_data = normalized_student_data[300:350]
28 student_test_labels = rank_one_hot[300:350]
29
30 student_val_data = normalized_student_data[350:]
31 student_val_labels = rank_one_hot[350:]
```

executed in 12ms, finished 13:45:11 2019-02-19

(400, 2)

TO DO: Define and train a network, then plot the accuracy of the training, validation, and testing

1. Use a validation set
2. Propose and train a network
3. Print the history of the training
4. Evaluate with a test set

In [24]:

```
1 student_model = models.Sequential()
2 student_dropout = 0.3
3
4 student_model.add(layers.Dense(128, activation= 'sigmoid', input_shape=(2,)))
5
6 student_model.add(layers.Dropout(student_dropout))
7
8 student_model.add(layers.Dense(10, activation= 'sigmoid'))
9 student_model.add(layers.Dense(5, activation= 'sigmoid'))
10
11 student_model.compile(
12     optimizer= "rmsprop",
13     loss= "binary_crossentropy",
14     metrics=[ "accuracy" ]
15 )
16
17 student_model.summary()
```

executed in 147ms, finished 13:50:34 2019-02-19

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_13 (Dense)	(None, 128)	384
dropout_5 (Dropout)	(None, 128)	0
dense_14 (Dense)	(None, 10)	1290
dense_15 (Dense)	(None, 5)	55
=====	=====	=====
Total params: 1,729		
Trainable params: 1,729		
Non-trainable params: 0		

In [25]:

```
1 student_model_history = student_model.fit(
2     student_train_data,
3     student_train_labels,
4     epochs= 16,
5     batch_size= 64,
6     validation_data= (student_val_data, student_val_labels),
7     verbose= 2
8 )
```

executed in 1.22s, finished 13:50:36 2019-02-19

Train on 300 samples, validate on 50 samples

Epoch 1/16

- 0s - loss: 0.6999 - acc: 0.6373 - val_loss: 0.6536 - val_acc: 0.6960

Epoch 2/16

- 0s - loss: 0.6529 - acc: 0.6487 - val_loss: 0.6241 - val_acc: 0.6960

Epoch 3/16

- 0s - loss: 0.6258 - acc: 0.6507 - val_loss: 0.6031 - val_acc: 0.6960

Epoch 4/16

- 0s - loss: 0.6048 - acc: 0.6580 - val_loss: 0.5872 - val_acc: 0.6960

Epoch 5/16

- 0s - loss: 0.5900 - acc: 0.6647 - val_loss: 0.5744 - val_acc: 0.6960

Epoch 6/16

- 0s - loss: 0.5773 - acc: 0.6640 - val_loss: 0.5638 - val_acc: 0.6960

Epoch 7/16

- 0s - loss: 0.5650 - acc: 0.6713 - val_loss: 0.5548 - val_acc: 0.6960

Epoch 8/16

- 0s - loss: 0.5552 - acc: 0.7013 - val_loss: 0.5470 - val_acc: 0.7640

Epoch 9/16

- 0s - loss: 0.5478 - acc: 0.7560 - val_loss: 0.5402 - val_acc: 0.8000

Epoch 10/16

- 0s - loss: 0.5401 - acc: 0.7933 - val_loss: 0.5339 - val_acc: 0.8000

Epoch 11/16

- 0s - loss: 0.5335 - acc: 0.7993 - val_loss: 0.5282 - val_acc: 0.8000

Epoch 12/16

- 0s - loss: 0.5268 - acc: 0.7993 - val_loss: 0.5229 - val_acc: 0.8000

Epoch 13/16

- 0s - loss: 0.5213 - acc: 0.8000 - val_loss: 0.5179 - val_acc: 0.8000

Epoch 14/16

- 0s - loss: 0.5167 - acc: 0.8000 - val_loss: 0.5134 - val_acc: 0.8000

Epoch 15/16

- 0s - loss: 0.5114 - acc: 0.8000 - val_loss: 0.5091 - val_acc: 0.8000

Epoch 16/16

- 0s - loss: 0.5072 - acc: 0.8000 - val_loss: 0.5050 - val_acc: 0.8000

In [26]:

```
1 student_result = student_model.evaluate(  
2     student_test_data,  
3     student_test_labels  
4 )  
5  
6 print('Student score: {}'.format(student_result[1] * 100))
```

executed in 12ms, finished 13:50:41 2019-02-19

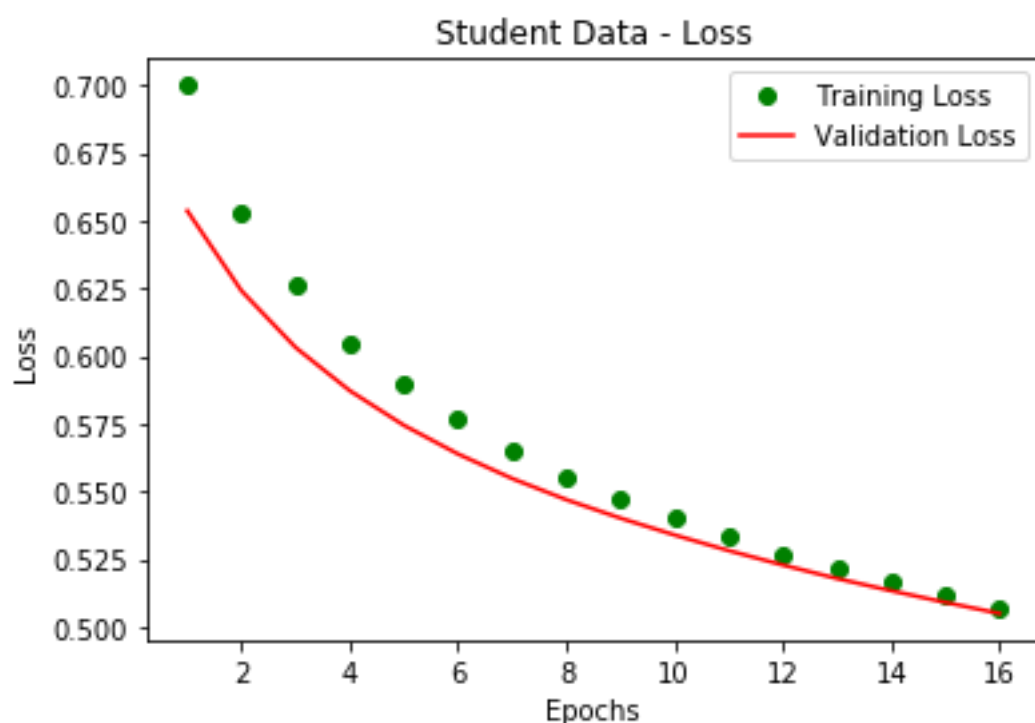
50/50 [=====] - 0s 132us/step

Student score: 80.00000500679016

In [27]:

```
1 student_history = student_model_history.history  
2  
3 student_val_loss = student_history['val_loss']  
4 student_loss     = student_history['loss']  
5 student_epochs   = range(1, len(student_loss) + 1)  
6  
7 plt.plot(student_epochs, student_loss, 'go', label= "Training Loss")  
8 plt.plot(student_epochs, student_val_loss, 'r', label= "Validation Loss")  
9 plt.title("Student Data - Loss")  
10 plt.ylabel("Loss")  
11 plt.xlabel("Epochs")  
12 plt.legend()  
13 plt.show()
```


executed in 193ms, finished 13:50:43 2019-02-19



In [28]:

```
1 student_val_acc = student_history['val_acc']
2 student_acc      = student_history['acc']
3
4 plt.plot(student_epochs, student_acc, 'go', label= "Training Acc")
5 plt.plot(student_epochs, student_val_acc, 'r', label= "Validation Acc")
6 plt.title("Student Data - Accuracy")
7 plt.ylabel("Accuracy")
8 plt.xlabel("Epochs")
9 plt.legend()
10 plt.show()
```

executed in 197ms, finished 13:50:44 2019-02-19



Student conclusion

This model was the hardest given the small dataset provided, we had 400 rows for training, validation & testing.

I believe that my preprocessing work flow could be better due that I do not use any statistical function to better the data.

In []:

1

